

HIERARCHICAL FAILURE CONTAINMENT: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study hierarchical robot policies. The motivating bet is: Design skill hierarchies that contain low-level failures before they corrupt task state.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For hierarchical robot policies, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026; 2025a;b;c;d; 2011). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.532	0.468
observed_only_wam	hidden_mode_shift	0.168	0.833
augmentation_as_data_wam	nominal	0.568	0.432
augmentation_as_data_wam	hidden_mode_shift	0.268	0.732
scalar_uncertainty_wam	nominal	0.620	0.380
scalar_uncertainty_wam	hidden_mode_shift	0.357	0.643
proposed_branch_mechanism	nominal	0.792	0.208
proposed_branch_mechanism	hidden_mode_shift	0.537	0.463

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Hierarchical Failure Containment argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- State parametrization in optimal control problems and low level controller design for multiple autonomous mobile robot systems, 2011. <https://doi.org/10.1109/aim.2011.6027135>.
- Grounding large language models for robot task planning using closed-loop state feedback, 2025a. <https://doi.org/10.1002/adrr.202500072>.
- Hierarchical vision-language planning for multi-step humanoid manipulation, 2025b. <http://arxiv.org/abs/2506.22827v3>.
- Hierarchical vision language action model using success and failure demonstrations, 2025c. <http://arxiv.org/abs/2512.03913v1>.
- Failsafe: Reasoning and recovery from failures in vision-language-action models, 2025d. <http://arxiv.org/abs/2510.01642v2>.
- H-wm: Robotic task and motion planning guided by hierarchical world model, 2026. <http://arxiv.org/abs/2602.11291v2>.